

# Analysis of Levenshtein Distance for DNA Sequences

Shivek

December 1, 2025

The **Levenshtein distance** is a "string metric" that measures the difference between two sequences. It is formally defined as the **minimum number of single-character edits** (insertions, deletions, or substitutions) required to change one string into the other.

This algorithm is a foundational concept in bioinformatics, forming the basis for more complex sequence alignment algorithms.

---

## 1 Basic Algorithm

The Levenshtein distance is calculated using a **dynamic programming** approach. The core idea is to build a 2D matrix (or grid) that stores the distances between all prefixes of the two input strings.

Let's say we have two strings,  $A$  of length  $m$  and  $B$  of length  $n$ . We create a matrix  $D$  of size  $(m + 1) \times (n + 1)$ .

A cell  $D[i][j]$  in this matrix represents the Levenshtein distance between the first  $i$  characters of string  $A$  (i.e.,  $A[1\dots i]$ ) and the first  $j$  characters of string  $B$  (i.e.,  $B[1\dots j]$ ).

The algorithm fills this matrix using two steps:

### Step 1: Initialization

The first row and first column are the "base cases."

- $D[i][0] = i$ : The distance from a string of length  $i$  to an empty string is  $i$  deletions.
- $D[0][j] = j$ : The distance from an empty string to a string of length  $j$  is  $j$  insertions.

### Step 2: Recurrence Relation

We fill the rest of the matrix, cell by cell. For any cell  $D[i][j]$ , we have three choices to get to this state:

1. **Deletion:** We had the solution for  $D[i - 1][j]$  (aligning  $A[1\dots i - 1]$  with  $B[1\dots j]$ ) and we **delete**  $A[i]$ . Cost:  $D[i - 1][j] + 1$ .
2. **Insertion:** We had the solution for  $D[i][j - 1]$  (aligning  $A[1\dots i]$  with  $B[1\dots j - 1]$ ) and we **insert**  $B[j]$ . Cost:  $D[i][j - 1] + 1$ .
3. **Substitution (or Match):** We had the solution for  $D[i - 1][j - 1]$  (aligning  $A[1\dots i - 1]$  with  $B[1\dots j - 1]$ ). We now look at  $A[i]$  and  $B[j]$ .

- If  $A[i] == B[j]$  (it's a **match**), no cost is added. Cost:  $D[i - 1][j - 1] + 0$ .
- If  $A[i] \neq B[j]$  (it's a **substitution**), we add 1. Cost:  $D[i - 1][j - 1] + 1$ .

The algorithm is "greedy" and always takes the minimum of these three possibilities. The complete recurrence relation is:

$$D[i][j] = \min \begin{cases} D[i - 1][j] + 1 & \text{(Deletion from A)} \\ D[i][j - 1] + 1 & \text{(Insertion into A)} \\ D[i - 1][j - 1] + \text{cost} & \text{(Match/Substitution)} \end{cases}$$

Where cost is 0 if  $A[i] == B[j]$  and 1 otherwise.

The final answer (the Levenshtein distance between all of  $A$  and all of  $B$ ) is the value in the bottom-right cell:  $D[m][n]$ .

---

## 2 Time and Space Complexity Analysis

### 2.1 Time Complexity: $O(m \times n)$

- The algorithm must fill every cell in the  $(m + 1) \times (n + 1)$  matrix.
- The calculation for each cell,  $D[i][j]$ , involves three lookups, two additions, and one min operation, all of which are constant time,  $O(1)$ .
- Therefore, the total time complexity is the number of cells multiplied by the constant work per cell, which is  $O(m \times n)$ .

### 2.2 Space Complexity: $O(m \times n)$

- The default algorithm requires storing the entire  $(m + 1) \times (n + 1)$  matrix in memory to compute the final value.
- This results in a space complexity of  $O(m \times n)$ .

**Optimization:** This can be optimized to  $O(\min(m, n))$  space. To compute the *current* row  $i$ , we only need the values from the *previous* row  $i - 1$ . We never need to look back further. By storing only two rows (the "previous" and "current") and alternating between them, we can reduce the space complexity significantly.

---

## 3 Proof of Correctness (Why it Works)

The algorithm's correctness is proven by induction.

1. **Base Case:** The initialization (Step 1) is correct. The edit distance from an empty string to a string of length  $j$  is  $j$  insertions (e.g., " $\text{"}$ "  $\rightarrow$  " $\text{ATG}$ " is 3 insertions).  $D[0][j] = j$  is correct. The same logic applies to  $D[i][0]$ .
2. **Inductive Hypothesis:** Assume that for all  $i' < i$  and  $j' < j$ , the sub-problems  $D[i'][j']$  have been solved correctly and store the true Levenshtein distance.

3. **Inductive Step:** We must prove that  $D[i][j]$  is calculated correctly. Any optimal sequence of edits that transforms  $A[1\dots i]$  to  $B[1\dots j]$  *must* end in one of three operations:

- **Case 1 (Deletion):** The last operation is deleting  $A[i]$ . This means we first optimally transformed  $A[1\dots i-1]$  to  $B[1\dots j]$ . The total cost is (cost of that optimal transform) + 1. By our hypothesis, this is  $D[i-1][j] + 1$ .
- **Case 2 (Insertion):** The last operation is inserting  $B[j]$ . This means we first optimally transformed  $A[1\dots i]$  to  $B[1\dots j-1]$ . The total cost is  $D[i][j-1] + 1$ .
- **Case 3 (Substitution/Match):** The last operation involved  $A[i]$  and  $B[j]$ . This means we first optimally transformed  $A[1\dots i-1]$  to  $B[1\dots j-1]$ .
  - If  $A[i] == B[j]$ , no new cost is added. Total cost:  $D[i-1][j-1] + 0$ .
  - If  $A[i] \neq B[j]$ , the last operation was a substitution. Total cost:  $D[i-1][j-1] + 1$ .

Since the Levenshtein distance is the *minimum* number of edits,  $D[i][j]$  must be the minimum of these three possible cases. This matches the recurrence relation. By induction, the entire matrix is filled correctly, and  $D[m][n]$  holds the final answer.

---

## 4 Pseudo Code

Note: This pseudo-code assumes 0-based indexing for strings (e.g.,  $A[0]$  is the first character) and 1-based indexing for the DP matrix loops.

```
function LevenshteinDistance(string A, string B):
    // m = length of A, n = length of B
    m = A.length
    n = B.length

    // 1. Initialize the (m+1) x (n+1) matrix D
    D = new int[m + 1][n + 1]

    // 2. Fill the base cases (first row and column)
    for i from 0 to m:
        D[i][0] = i // Cost of deleting i chars from A to get ""

    for j from 0 to n:
        D[0][j] = j // Cost of inserting j chars to "" to get B

    // 3. Fill the rest of the matrix
    for i from 1 to m:
        for j from 1 to n:

            // Check if characters at A[i-1] and B[j-1] are the same
            // (We use i-1 and j-1 because strings are 0-indexed)
            if A[i-1] == B[j-1]:
                cost = 0
            else:
```

```

cost = 1

// Calculate costs for each operation
deletion = D[i-1][j] + 1
insertion = D[i][j-1] + 1
substitution = D[i-1][j-1] + cost

// The cell D[i][j] gets the minimum of these three
D[i][j] = min(deletion, insertion, substitution)

// 4. Return the final answer
return D[m][n]

```

---

## 5 Dry Run (Biological Sequences)

Let's find the Levenshtein distance between two short DNA sequences:

- $A = \text{"CAT"}$  ( $m = 3$ )
- $B = \text{"TAG"}$  ( $n = 3$ )

We create a  $(3 + 1) \times (3 + 1) = 4 \times 4$  matrix.

Table 1: Dry Run Matrix for "CAT" vs "TAG"

		$B \rightarrow$			
		(empty)	T	A	G
$A \downarrow$	D	j = 0	j = 1	j = 2	j = 3
(empty)	i = 0	0	1	2	3
C	i = 1	1	1	2	3
A	i = 2	2	2	1	2
T	i = 3	3	2	2	2

**Final Answer:** The Levenshtein distance is **2**.

This makes sense:  $C \rightarrow T$  (substitution),  $A \rightarrow A$  (match),  $T \rightarrow G$  (substitution). Total: 2 substitutions.

---

## 6 Graphical Visualization

The completed dry run table *is* the graphical visualization. We can also visualize the **edit path** by backtracking from the final cell  $D[3][3]$  to  $D[0][0]$ , always moving to the cell that generated the minimum value.

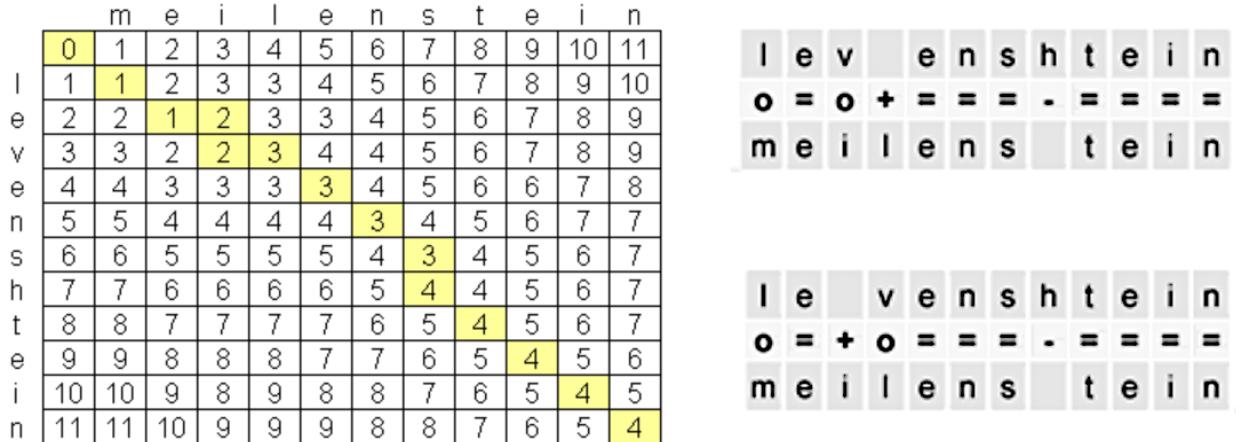


Figure 1: Visualization of the DP matrix

Table 2: Matrix with Backtrace Path ( $\leftarrow$ )

$A \downarrow$	D	$B \rightarrow$			
		(empty)	$j = 0$	$j = 1$	$j = 2$
(empty)	$i = 0$	0	$0 \leftarrow$	$1 \leftarrow$	2
C	$i = 1$	1	1	$1 \leftarrow$	2
A	$i = 2$	2	2	$1 \leftarrow$	2
T	$i = 3$	3	2	2	2

#### Backtrace Path:

1. Start at  $D[3][3]$  (value 2): Came from  $D[2][2]$  (a substitution, since  $T \neq G$ ).
2. At  $D[2][2]$  (value 1): Came from  $D[1][1]$  (a match, since  $A == A$ ).
3. At  $D[1][1]$  (value 1): Came from  $D[0][0]$  (a substitution, since  $C \neq T$ ).

This traceback gives us the full alignment:

C A T  
 T A G  
 (Sub, Match, Sub)  $\rightarrow$  2 edits.

## 7 Situational Performance (On DNA Sequences)

Levenshtein distance is a *good* algorithm, but it's often *not* the best tool for biological sequences.

## 7.1 Strengths

1. **Simplicity and Accuracy:** It's a simple, easy-to-implement algorithm that gives a mathematically correct *edit distance*.
2. **Good for Short, Similar Sequences:** It works very well for comparing two short reads, or checking for SNPs (Single Nucleotide Polymorphisms) between two highly similar sequences.

## 7.2 Weaknesses (Major Performance Issues for DNA)

1. **Biologically Naive:** This is the biggest problem. Levenshtein treats all operations as equal (cost=1). In biology, this is wrong:
  - **Substitutions:** A substitution of  $A \rightarrow G$  (a **transition**) is *far* more common and biologically likely than  $A \rightarrow T$  (a **transversion**). Levenshtein can't do this.
  - **Gaps:** An insertion or deletion (an "indel" or "gap") is often a single biological event, but it might delete 10 bases at once. Levenshtein would count this as 10 separate edits. Proper alignment algorithms use an "affine gap penalty".
2. **It's a Global Alignment:** The standard Levenshtein algorithm forces the *entire* string  $A$  to match the *entire* string  $B$ . This is almost *never* what we want in bioinformatics.
  - **Problem:** We usually want to find a *small* gene (e.g., 1,000 bases) inside a *huge* chromosome (e.g., 250,000,000 bases).
  - **Solution:** We need **local alignment**.
3. **Poor Time/Space Complexity for Genomes:** The  $O(m \times n)$  complexity is completely unusable for large-scale genomics.
  - Comparing two human chromosomes ( $m \approx 250M, n \approx 250M$ ) would require a matrix with  $6.25 \times 10^{16}$  cells. This is not computationally feasible.

**Conclusion:** For *actual performance on DNA sequences*, it is a poor choice.

## 8 Adaptation for Fuzzy Pattern Matching in DNA

The standard algorithm (Section 1) is a **global alignment**, comparing the *entire* string  $A$  to the *entire* string  $B$ . This is useless for finding a small pattern (e.g., a 6-base-pair motif) inside a large genome (e.g., 4.6 million base pairs).

To find all substrings in a large Text  $T$  that "fuzzily" match a Pattern  $P$  with at most  $k$  edits, we must adapt the algorithm.

Let  $P$  be the pattern of length  $m$  (on the vertical axis) and  $T$  be the genome text of length  $n$  (on the horizontal axis).

### 8.1 The "Slight Modification"

The algorithm is modified in two ways: the initialization and the final result-finding.

### 8.1.1 Modified Initialization

The key is to allow a match to **start at any point** in the text  $T$  without penalty. We achieve this by setting the first row to all zeros.

- $D[0][j] = 0$  for all  $j$  from 0 to  $n$  (text length): This means the cost of matching an empty pattern to *any* prefix of the text is 0. This "zeros" the cost for starting a new match.
- $D[i][0] = i$  for all  $i$  from 0 to  $m$  (pattern length): This remains the same. The cost to match a pattern of length  $i$  to an empty string is  $i$  deletions.

### 8.1.2 Modified Result Finding

The recurrence relation (Section 1.2) is filled exactly the same way. However, the answer is no longer in the single cell  $D[m][n]$ .

Instead, we must scan the **entire last row** (row  $m$ ).

- A value  $D[m][j]$  in the last row represents the minimum Levenshtein distance required to match the *entire* pattern  $P$  to a substring of the text  $T$  that *ends* at position  $j$ .
- **The Result:** We iterate  $j$  from 1 to  $n$ . If  $D[m][j] \leq k$  (where  $k$  is our error threshold), we record  $j$  as the *end position* of a fuzzy match.

## 8.2 Complexity of the Adaptation

The time and space complexity remain the same as the standard algorithm.

- **Time:**  $O(m \times n)$ . This is why it is linear  $O(n)$  for a fixed pattern  $m$ , but slow  $O(m)$  for a fixed text  $n$  and growing pattern  $m$ .
  - **Space:**  $O(n)$  with the standard row-optimization.
-