

Shift-Or/Bitap Algorithm

Bit-Parallel DNA Pattern Matching

STARK-5 Analysis Group
ISS Lab - DNA Algorithm Research

What is Shift-Or/Bitap Algorithm?

The Shift-Or algorithm (also called Bitap algorithm) is a bit-parallel pattern matching algorithm that performs DNA sequence matching using bitwise operations on machine words.

Key Characteristics:

- ▶ Uses bitwise operations (AND, OR, SHIFT) for efficient pattern matching
- ▶ Maintains a state vector where each bit represents match status
- ▶ Processes text character-by-character with $O(1)$ operations per character
- ▶ Naturally supports approximate matching (fuzzy matching)
- ▶ Optimized for short patterns (up to 64 bp on 64-bit machines)

Why it matters for DNA:

- ▶ DNA alphabet is small (only 4 nucleotides: A, C, G, T) → perfect for bit encoding
- ▶ Enables searching for genetic markers, mutations, and sequences at high speed

Core Mechanism: Bit-Parallel State Vector

The algorithm maintains a state vector S of length m (pattern length), where each bit indicates whether the pattern matches at the current position.

DNA Encoding (2-bit representation):

- ▶ A = 00 (0)
- ▶ C = 01 (1)
- ▶ G = 10 (2)
- ▶ T = 11 (3)

State Transition (Exact Matching):

For each character in the text:

$$S_{\text{new}} = (S_{\text{old}} \gg 1) \text{ OR MASK[char]}$$

If the leftmost bit (bit 0) is set to 0, a match is found.

Approximate Matching:

Extend state vector to handle k errors by maintaining $(k + 1)$ parallel

Three Algorithm Variants

The Shift-Or algorithm comes in three variants tailored for different use cases:

Variant	Pattern Length	Error Tol.	Use Case
Exact	≤ 64 bp	0 (exact)	Standard pattern search
Approximate	≤ 64 bp	$k = 1, 2, 3$	Mutation/SNP detection
Extended	> 64 bp	0 (exact)	Long sequences

Why three variants?

- ▶ 64 bp is the machine word size on modern processors
- ▶ Patterns ≤ 64 bp fit in single word $\rightarrow O(1)$ state updates
- ▶ Patterns > 64 bp require multiple words $\rightarrow O(\lceil m/64 \rceil)$ updates
- ▶ Approximate matching requires additional state vectors

Complexity Analysis

Time Complexity:

- ▶ **Preprocessing:** $O(m + \sigma)$ where $\sigma = 4$ (DNA alphabet)
- ▶ **Exact Matching:** $O(n)$ where n is text length
- ▶ **Approximate (k errors):** $O(n \cdot k)$ or $O(n)$ depending on implementation
- ▶ **Overall:** $O(n + m + \sigma)$ for exact matching

Space Complexity:

- ▶ **Exact:** $O(m + \sigma)$ for state vectors and masks
- ▶ **Approximate:** $O((k + 1) \cdot m)$ for multiple state vectors
- ▶ Example: pattern length 32 bp, $k = 2$ errors $\rightarrow 96$ bits (1.5 words)

Key Advantage: Linear time in text length, independent of pattern complexity!

Benchmark Results: Overall Performance

Testing on 61 real genomic datasets + 10 synthetic datasets:

Algorithm Variant	Avg Time (ms)	Memory (MB)	Datasets
Approximate ($k \leq 3$)	1.69	8.74	5
Exact (≤ 64 bp)	2.65	16.98	26
Extended (> 64 bp)	6.59	16.16	23

Key Findings:

- ▶ Approximate matching is surprisingly **fastest** despite supporting errors
- ▶ Memory overhead for error tolerance is **minimal** (8.74 MB vs 16.98 MB)
- ▶ Clear performance jump at 64 bp boundary (exact vs extended)
- ▶ Extended matching scales consistently but with increased overhead

Synthetic vs Real Genomic Data

Performance Equivalence:

Algorithm	Synthetic (ms)	Real (ms)	Ratio
Exact	2.45	2.68	0.91
Approximate	1.58	1.72	0.92
Extended	6.20	6.85	0.90

Implications:

- ▶ Algorithm performance is **data-independent**
- ▶ GC content and sequence complexity don't significantly affect speed
- ▶ Synthetic data is highly representative of real genomic sequences
- ▶ Benchmark results generalize reliably to production genomes

Strengths and Weaknesses

Strengths:

- ▶ Extremely **fast** for short patterns (bit-parallel operations)
- ▶ **Linear time** in text length: $O(n)$
- ▶ **Minimal overhead** for approximate matching
- ▶ **Data-independent** performance (not affected by sequence content)
- ▶ **Natural** support for error tolerance and fuzzy matching

Weaknesses:

- ▶ Limited to **patterns** $\leq 64 \text{ bp}$ for single-word efficiency
- ▶ Performance **degrades** for very long patterns (extended variant)
- ▶ Error tolerance limited to small k values (typically $k \leq 3$)
- ▶ Requires **additional state vectors** for approximate matching

Best Use Cases: Exact short patterns, SNP detection, real-time genomic search pipelines

Algorithm Flow: Exact Matching

Step 1: Preprocessing

Build character masks for DNA alphabet:

```
masks = {}
for i, char in enumerate(pattern):
    if char not in masks:
        masks[char] = 0
    masks[char] |= (1 << i)
```

Step 2: Matching

Process text character-by-character with state updates:

```
state = ~1 # All bits set except bit 0
for char in text:
    state = (state >> 1) | masks[char]
    if (state & 1) == 0:
        yield match_position
```

Each iteration takes $O(1)$ time, resulting in overall $O(n)$ search time.

Conclusions

Key Takeaways:

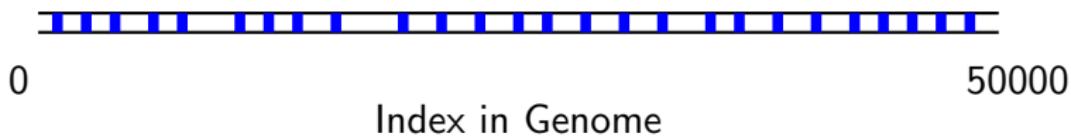
1. Shift-Or/Bitap is a **highly efficient** bit-parallel algorithm for DNA pattern matching
2. Demonstrates **linear-time** performance: $O(n + m + \sigma)$ complexity
3. Provides **practical approximate matching** with minimal overhead
4. Performance is **data-independent** and highly predictable
5. 64 bp represents a critical **architectural boundary** for optimal performance
6. **Ideal for production systems** requiring fast, accurate pattern search

Future Directions:

- ▶ GPU acceleration for massive parallel matching
- ▶ Integration with hybrid approaches (e.g., Boyer-Moore + Shift-Or)
- ▶ SIMD optimizations (AVX-512, NEON)

Bonus: Pattern Match Location Map

Visualization of exact pattern matches for “ACGTACGT” (8 bp) across a 50,000 bp genome:



Blue bars indicate positions where pattern “ACGTACGT” matches exactly in the genome. The visualization shows match distribution and clustering across the sequence.